

INTRODUCTORY GUIDE TO MESSAGE PASSING IN DISTRIBUTED SYSTEMS

Florian Willich
Hochschule für Technik und Wirtschaft Berlin
University of Applied Sciences Berlin
Course: Distributed Systems
Lecturer: Prof. Dr. Christin Schmidt

June 14, 2015

Abstract

Message passing in distributed systems is a model to exchange messages within a process pair by making use of several standards and implementation details. Those have been developed to offer the right message passing models for the different areas of applications. The programming language Erlang natively supports an asynchronous message passing model which makes the implementation of concurrent applications transparent to the software developer.

Contents

1	MESSAGE PASSING	3
1.1	INTRODUCTION	3
1.2	REQUIREMENTS IN THEORY	3
1.3	REQUIREMENTS PROVIDER IN PRACTICE	4
1.4	ASYNCHRONOUS VS. SYNCHRONOUS MESSAGE PASSING	4
2	MESSAGE-PASSING INTERFACE STANDARD	5
3	DEFINITION	6
4	MESSAGE PASSING IN ERLANG	6
4.1	INTRODUCTION	6
4.2	CONCURRENCY IN ERLANG	7
4.3	IMPLEMENTATION OF A SIMPLE DIFFIE-HELLMAN KEY EXCHANGE ALGORITHM	8
4.4	RUNNING THE APPLICATION	12
4.4.1	Preparatory Work	12

4.4.2	On One Node	12
4.4.3	On Two Nodes	13
5	CLOSING REMARKS	13
A	my_math Erlang Module	14
	References	16

1 MESSAGE PASSING

1.1 INTRODUCTION

To provide services and execute tasks, a distributed system has to have a proficient communication model implementation. There are several models of communication in distributed systems such as Remote Procedure Calls (RPC) or object method invocation (Tanenbaum & Steen, 2007, chap. 4.3 / p. 203). With this paper I'm giving an introduction to the communication models of *message passing*.

As human beings we already have a deep understanding of different models of message passing. While you read these words, I am passing a message to you, which seems to be a trivial thing to people who can read. In more philosophical terms, now that you are reading those words, a message is being passed from one entity (me, the writer) to another entity (you, the reader) which is obviously non trivial considering all the assumptions we would have to make to actually realize this message passing model between human beings.

Message passing in distributed systems is based on messages, composed of bit strings, exchanged within a process pair which would be the equivalent to the entity pair. It is important to understand that message passing is a model designed for Inter-Process Communication (IPC). Whether those processes are located on one or on two systems, is irrelevant to the provided functionality.

1.2 REQUIREMENTS IN THEORY

The following items are the theoretical basic requirements for a message passing model:

- **Connectivity:** A connection to communicate has to be established between the process pair
- **Ability:** Each process has to be able to receive or send messages
- **Integrity:** Sent messages have to be delivered as is
- **Intelligibility:** The receiving process has to be able to interpret the message as intended

I intentionally left out the requirement of executability without there is no effect when the message is passed. To call the desired instructions in the receiving process is an implementation detail of executed code and thus not a requirement of the message passing model.

1.3 REQUIREMENTS PROVIDER IN PRACTICE

To provide the above mentioned requirements several message passing models resulting in well defined standards and concrete implementations have been developed in the last decades. To facilitate understanding the relations between the theoretical requirements and a concrete implementation, I chose the Transmission Control Protocol / Internet Protocol (TCP/IP) which uses several socket primitives as example (Tanenbaum & Steen, 2007, chap 4.3.1 on p. 141 - 142):

- **Connectivity:** TCP/IP provides the *Socket* primitive creating a socket end point and also the *Bind* primitive to bind a local address to that socket. The *Connect* primitive then provides the functionality to establish a connection.
- **Ability:** TCP/IP provides the two primitives *Send* and *Receive* to simply send and receive data via the connection.
- **Integrity:** TCP/IP provides several mechanisms to ensure that there has been no data loss when sending or receiving messages. On the other hand this makes the protocol slower than other protocols such as the Real-Time Transport Protocol (RTP) or the User Datagram Protocol (UDP).

The requirements of **Intelligibility** do not fall within the responsibility of TCP/IP and thus other standards have to take place, such as Extensible Markup Language (XML) or JavaScript Object Notation (JSON) to structure the messages in a standardized manner.

There are other socket primitives than the above mentioned to provide the functionality of TCP/IP (Tanenbaum & Steen, 2007, ch. 4.3.1 on p. 141). The referenced book (Tanenbaum & Steen, 2007) provides useful and valuable information on this topic.

It is also important to understand that this was just a round-up of how one can make use of the socket primitives to implement message passing in an application. There is plenty to discuss about how message passing is realized, such as the Open Systems Interconnection Reference Model (OSI) which was developed by the International Organization for Standardization (ISO) to model the different layers of network oriented communication (Tanenbaum & Steen, 2007, ch. 4.1.1 on p. 116). Furthermore, an important detail is that the Operating System (OS) has to reserve local memory to provide a buffer for the in- and outgoing messages.

1.4 ASYNCHRONOUS VS. SYNCHRONOUS MESSAGE PASSING

After discussing a protocol to handle the socket primitives, it is necessary to consider one major difference in message passing models, which is whether to

use asynchronous or synchronous message passing. Passing a message synchronously means that the called *message send routine* returns after the request has been successfully transmitted (or an error occurred). Receiving a message synchronously means that the called message receiving routine reads a specific amount of bytes from the socket and returns that message.

To pass or receive a message asynchronously, an additional application layer i.e. middleware is introduced. The call of a send routine can either return when the middleware took over the transmission of the request, or when it successfully sent the request to the receiver, or when it successfully sent the request to the receiver, assuring this by a corresponding message (Tanenbaum & Steen, 2007, ch. 4.1 on p. 125).

The middleware can also perform some preparatory work e.g. separating stand-alone request bit strings and parsing them into a data structure and returning it when calling the receive routine of the message passing middleware.

One key model to provide such middleware facilities are message queues. A message queue is a queue to store messages by the First In - First Out (FIFO) principle. This enables the application to asynchronously send and receive messages by offering an incoming message queue and an outgoing message queue. This makes all the send and receive mechanisms fully transparent to the application (Tanenbaum & Steen, 2007, ch. 4.3.2 on p. 145 - 147).

2 MESSAGE-PASSING INTERFACE STANDARD

To pass messages between processes is adding large overhead to the communication model. Nevertheless, providing a well defined interface to control exactly how messages are passed regains the ability to write High-Performance Computing (HPC) applications.

The Message-Passing Interface (MPI) is a message-passing library interface specification, created for high performance and scalable distributed systems where high-efficiency is needed. The MPI standard was designed by the Message-Passing Interface Forum (MPIF) which is an open group with representatives from many organizations. The current version of the MPI standard is MPI-3.0 (Message Passing Interface Forum, 2012, ch. Abstract/ii & ch. Acknowledgements/xx & ch. 1.1 on p. 1 & ch. 1.2 on p. 2).

The MPIF aims to offer a standard which establishes a practical, portable, efficient and flexible way of implementing message-passing in various high-level programming languages (e.g. C, C++, Fortran) (Message Passing Interface Forum, 2012, ch. History/iii). Furthermore, the MPI simplifies the communication primitives and brings them to an abstraction level to perfectly fit the programmer's needs of writing efficient and clean code for such HPC

distributed systems (Tanenbaum & Steen, 2007, ch. 4.3.1 on p 143).

TCP/IP does not fit those requirements. While the socket primitives *read* and *write* are sufficient for several general-purpose protocols (managing the communication across networks), they are insufficient for high-speed interconnection networks such as super computers or server clusters. The MPI standard offers a set of functions and datatypes with which the software developer is able to explicitly execute synchronous and asynchronous message passing routines (Tanenbaum & Steen, 2007, ch. 4.3 on p. 143).

The MPIF offers a very detailed technical report (Message Passing Interface Forum, 2012) of the MPI standard which is not only defining the standard but offers detailed information on the organization and motivation.

3 DEFINITION

The given task of Prof. Dr. Christin Schmidt was to write a technical report defining the term *Message Passing*. After discussing the theoretical basis and practical implementations, message passing in distributed systems appears to be a technical jargon term to characterize a communication model in distributed systems that deals with messages:

Message passing in distributed systems is a model to exchange messages within a process pair by making use of several standards and implementation details. The specifically used message passing model can diverge extremely in its provided facilities and defines how to establish the connection, send and receive messages.

However, no message passing model defines the medium that transports the messages nor the outcome of a message and thus the following delimitation has to be made:

The physical conditions and the implementation of executing the desired instructions is not defined by any message passing model.

4 MESSAGE PASSING IN ERLANG

The above written discussions always led to the conclusion that message passing is not only adding computational overhead but also adds a lot of things the programmer has to keep track of. The message passing model implemented in Erlang simplifies and abstracts the use of message passing so the programmer can again focus on the application logic.

4.1 INTRODUCTION

Erlang is a functional, declarative programming language that was written for the need of real-time, non-stop, concurrent, very large and distributed system applications (Armstrong, 1996, chap. 1 / p. 1). While the language was originally designed by Joe Armstrong, Robert Virding and Mike Williams for Ericsson (a Swedish telecommunications provider) in 1986, it finally became open source in 1998, thanks to the open source initiatives lead by Linux (Däcker, 2000, chap. 8 on p. 39).

Joe Armstrong, who started to design Erlang by adding functionality to Prolog, named the new programming language after the Danish mathematician Agner Krarup Erlang (creator of the Erlang loss formula) following the tradition of naming programming languages after dead mathematicians (Däcker, 2000, chap. 4.1 on p. 13).

Erlang uses a native, asynchronous message passing model to communicate between light weight processes, also called actors (Armstrong, 1996, chap. 1 on p. 1). Erlang does not make heavy use of the executing OS to provide the described concurrency model, which implicitly means that Erlang decouples the underlying OS and thus providing a cross-platform and transparent message passing model (Armstrong, 1996, chap. 1 & 3 on p. 1 - 3).

4.2 CONCURRENCY IN ERLANG

Erlang provides semantics and built-in functions to parallelize applications by message passing (Ericsson AB, 2015, ch. 4.3 on p. 95 - 104):

- **spawn**(*Module*, *Exported Function*, *List of Arguments*)
Function which creates a new actor by running the *Exported Function* with the *List of Arguments* located in the *Module* (a set of functions located in one file) and which returns the Process Identifier (pid), which uniquely identifies the created actor.
- The **receive** construct allows the function executed by an actor to receive messages by using a message queue.
- The **!** operator sends the right-handed term to the left handed pid. The right-handed term is the sent message.
- **self()**
Function returning the pid of the actor executing the function.

Although the semantics and functions described above may require some fore-knowledge on the Erlang programming language, it was important for me to

show that Erlang is not only offering an easily intelligible interface for concurrency, it rather is a concurrent functional programming language. Nevertheless, there is plenty to discuss about the concurrency model of Erlang e.g. how to manage processes or handle errors. I recommend to take a look at the referenced official Erlang documentation or to *learn you some Erlang* on www.learnyousomeerlang.com.

4.3 IMPLEMENTATION OF A SIMPLE DIFFIE-HELLMAN KEY EXCHANGE ALGORITHM

The following implementation of a Diffie-Hellman key exchange algorithm in Erlang is simplified. Please note that this Erlang application is not in any way acceptable as an adequate implementation of the Diffie-Hellman key exchange algorithm for use in the field! The code demonstrates the simplicity of implementing concurrent applications in Erlang by using message passing. Due to the need for a pow function returning the datatype *Integer* I implemented my own pow function see Appendix A. The sourcecode is hosted by Github:

https://github.com/c-bebop/message_passing

Licensed under the MIT License. You're welcome to contribute!

```
1 %%% @author      Florian Willich
2 %%% @copyright   The MIT License (MIT) Copyright (c) 2014
3 %%%              University of Applied Sciences, Berlin, Germany
4 %%%              Florian Willich
5 %%%              For more detailed information, please read the
6 %%%              licence.txt in the erlang root directory.
7 %%% @doc         This module represents a simple implementation of
8 %%%              the Diffie-Hellman key exchange algorithm and is
9 %%%              part of my technical report 'Introductory Guide
10 %%%             to Message Passing in Distributed Systems'.
11 %%% @end
12 %%% Created 2015-06-06
13
14 -module(diffie_hellman).
15 -author("Florian Willich").
16 -export([ computeMyPublicKey/3,
17           computeSharedPrivateKey/3,
18           startKeyExchange/5,
19           listenKeyExchange/2,
20           startExample/0,
21           startRemoteExample/1,
22           startRemoteExample/5]).
23
24 %%% @doc Public data represents the data which is publicly shared
25 %%%      within two communication partners when exchanging keys
26 %%%      with the Diffie-Hellman key exchange algorithm.
27 %%%      p: The public prime number
28 %%%      g: The public prime number (1 ... p - 1)
29 %%%      componentKey: The computed component key
30 %%%      pid: The pid of the one who instantiated this record
31 %%%      name: The name of whom creates this public data.
```



```

32 -record(publicData, {p, g, componentKey, pid, name}).
33
34 %%% @doc Returns the value G to the power of MySecretKey Modulo P
35 %%% which is the public component key for the Diffie-Hellman
36 %%% key exchange algorithm.
37 %%% For more information see http://goo.gl/pzdiH
38 %%% @end
39 -spec computeMyPublicComponentKey(pos_integer(), pos_integer(),
    pos_integer()) -> pos_integer().
40 computeMyPublicComponentKey(P, G, MySecretKey) ->
41   my_math:pow(G, MySecretKey) rem P.
42
43 %%% @doc Returns the value ComponentKey to the power of
44 %%% MySecretKey Modulo P which is the private
45 %%% shared key for the Diffie-Hellman key exchange
46 %%% algorithm.
47 %%% For more information see http://goo.gl/pzdiH.
48 %%% @end
49 -spec computeSharedPrivateKey(pos_integer(), pos_integer(),
    pos_integer()) -> pos_integer().
50 computeSharedPrivateKey(P, ComponentKey, MySecretKey) ->
51   my_math:pow(ComponentKey, MySecretKey) rem P.
52
53 %%% @doc Starts the Diffie-Hellman key exchange algorithm by
54 %%% taking P (a prime number), G (1 ... P - 1), MySecretKey
55 %%% is the secret integer of the one who executes this
56 %%% function and the PartnerPID which is the PID of the
57 %%% communication partner with whom a key exchange shall be
58 %%% initiated. MyName shall be the name of the executing
59 %%% partner. This function sends the term
60 %%% {startKeyExchange, PublicData} to the PartnerPID where
61 %%% PublicData is of type publicData. Afterwards, the
62 %%% function starts a receive construct which is receiving
63 %%% the following:
64 %%% {componentKey, PublicData}:
65 %%% The message including all information needed for
66 %%% computing the private shared key and then prints it
67 %%% out.
68 %%% UnexpectedMessage:
69 %%% Prints out any unexpected incoming message and
70 %%% calls a recursion.
71 %%% After 3000 milliseconds:
72 %%% The function will return timeout.
73 %%% @end
74 -spec startKeyExchange(pos_integer(), pos_integer(), pos_integer(),
    term(), string()) -> term() | {error, atom()}.
75 startKeyExchange(P, G, MySecretKey, PartnerPID, MyName) ->
76   MyComponentKey = computeMyPublicComponentKey(P, G, MySecretKey),
77   MyPublicData = #publicData{p = P, g = G, componentKey =
    MyComponentKey, pid = self(), name = MyName},
78   PartnerPID ! {startKeyExchange, MyPublicData},
79
80   receive
81
82     {componentKey, #publicData{p = P, g = G, componentKey =
    PartnerComponentKey, pid = PartnerPID, name = PartnerName}}
    ->

```

```

83     PrivateSharedKey = computeSharedPrivateKey(P,
84         PartnerComponentKey, MySecretKey),
85     printSharedPrivateKey(self(), MyName, PartnerName,
86         PrivateSharedKey);
87
88     UnexpectedMessage ->
89     printUnexpectedMessage(UnexpectedMessage),
90     startKeyExchange(P, G, MySecretKey, PartnerPID, MyName)
91
92     after 3000 ->
93     {error, timeout_after_3000_ms}
94
95     end.
96
97     @doc Listens on Messages to start the Diffie-Hellman key
98     with the transferred MySecretKey and MyName
99     which shall be the name of the executing partner.
100     exchange by starting the following receive construct:
101     {startKeyExchange, PublicData}:
102     The message including all information needed to
103     start the key exchange by computing the own public
104     data which will then be send to the PartnerPID as
105     follows: {componentKey, MyPublicData}.
106     Afterwards, the private shared key will be printed
107     out and the function calls a recursion.
108     terminante:
109     Prints out that this function terminates with the
110     executing PID and returns ok.
111     UnexpectedMessage:
112     Prints out any unexpected incoming message and
113     calls a recursion.
114
115     @end
116
117     -spec listenKeyExchange(pos_integer(), string()) -> term().
118     listenKeyExchange(MySecretKey, MyName) ->
119     receive
120
121     {startKeyExchange, #publicData{p = P, g = G, componentKey =
122         PartnerComponentKey, pid = PartnerPID, name = PartnerName}}
123     ->
124
125     MyComponentKey = computeMyPublicComponentKey(P, G,
126         MySecretKey),
127     MyPublicData = #publicData{p = P, g = G, componentKey =
128         MyComponentKey, pid = self(), name = MyName},
129     PartnerPID ! {componentKey, MyPublicData},
130     PrivateSharedKey = computeSharedPrivateKey(P,
131         PartnerComponentKey, MySecretKey),
132     printSharedPrivateKey(self(), MyName, PartnerName,
133         PrivateSharedKey),
134     listenKeyExchange(MySecretKey, MyName);
135
136     terminate ->
137     io:format("~p terminates!~n", [self()]),
138     ok;
139
140     UnexpectedMessage ->
141     printUnexpectedMessage(UnexpectedMessage),
142     listenKeyExchange(MySecretKey, MyName)

```

```

132
133     end.
134
135 %%% @doc Prints out the UnexpectedMessage as follows:
136 %%% Received an unexpected message: 'Unexpected Message'
137 %%% @end
138 -spec printUnexpectedMessage(string()) -> term().
139 printUnexpectedMessage(UnexpectedMessage) ->
140     io:format("Received an unexpected message: ~p~n", [
141         UnexpectedMessage]).
142
143 %%% @doc Prints out the shared private key as follows:
144 %%% 'MyName' ('PID'): The shared private Key,
145 %%% exchanged with 'PartnerName' is: 'SharedKey'
146 %%% @end
147 -spec printSharedPrivateKey(term(), string(), string(), string())
148     -> term().
149 printSharedPrivateKey(PID, MyName, PartnerName, SharedKey) ->
150     io:format("~p (~p): The shared private Key, exchanged with ~p is:
151         ~p~n", [MyName, PID, PartnerName, SharedKey]).
152
153 %%% @doc Starts a key exchange example by spawning the Alice
154 %%% process, which executes the listenKeyExchange function
155 %%% with MySecretKey = 15, and the Bob process, which
156 %%% executes the startKeyExchange function with P = 23,
157 %%% G = 5, MySecretKey = 6 and PartnerPID = Alice. Returns
158 %%% {Alice, Bob} (pids).
159 %%%
160 -spec startExample() -> term().
161 startExample() ->
162     Alice = spawn(diffie_hellman, listenKeyExchange, [15, "Alice"]),
163     Bob = spawn(diffie_hellman, startKeyExchange, [23, 5, 6,
164         Alice, "Bob"]),
165     {Alice, Bob}.
166
167 %%% @doc Starts a key exchange remote example by spawning the
168 %%% Alice process, which executes the listenKeyExchange
169 %%% function with MySecretKey = 15, and the Bob process,
170 %%% located on the RemoteNode, which executes the
171 %%% startKeyExchange function with P = 23, G = 5,
172 %%% MySecretKey = 6 and Alice PID.
173 %%% Returns the pids of Alice and Bob.
174 %%% @end
175 -spec startRemoteExample(atom()) -> term().
176 startRemoteExample(RemoteNode) ->
177     Alice = spawn(RemoteNode, diffie_hellman, listenKeyExchange, [
178         15, "Alice"]),
179     Bob = spawn(diffie_hellman, startKeyExchange, [23, 5, 6,
180         Alice, "Bob"]),
181     {Alice, Bob}.
182
183 %%% @doc Starts a key exchange remote example by spawning the
184 %%% Alice process, which executes the listenKeyExchange
185 %%% function with AliceSecretKey, and the Bob process,
186 %%% located on the RemoteNode, which executes the
187 %%% startKeyExchange function with P, G,

```

```

182 %%%      BobSecretKey and Alice PID.
183 %%%      Returns the pids of Alice and Bob.
184 %%% @end
185 -spec startRemoteExample(atom(), pos_integer(), pos_integer(),
    pos_integer(), pos_integer()) -> term().
186 startRemoteExample(RemoteNode, P, G, BobSecretKey, AliceSecretKey)
    ->
187   Alice = spawn(RemoteNode, diffie_hellman, listenKeyExchange, [
    AliceSecretKey, "Alice"]),
188   Bob    = spawn(diffie_hellman, startKeyExchange, [P, G,
    BobSecretKey, Alice, "Bob"]),
189   {Alice, Bob}.

```

4.4 RUNNING THE APPLICATION

To run the applicatoin (4.3) in a Linux shell one shall do the following (assuming that Erlang is already installed, help can be found at http://www.erlang.org/doc/installation_guide/INSTALL.html).

4.4.1 Preparatory Work

Open a Linux shell, go to the *source* directory of the Erlang code and type:

```
1 $ erl -sname bob
```

Which will output (meta data depends of the executing system):

```

1 Erlang/OTP 17 [erts-6.3] [source] [64-bit] [smp:4:4] [async-threads
   :10] [hipe] [kernel-poll:false]
2
3 Eshell V6.3 (abort with ^G)
4 (bob@localhost)1>

```

The option *-sname bob* tells the Erlang shell to run on a node called *bob*. First one shall compile the two modules and afterwards:

```

1 (bob@localhost)1> c(diffie_hellman).
2 {ok,diffie_hellman}
3 (bob@localhost)2> c(my_math).
4 {ok,my_math}

```

4.4.2 On One Node

To run the application on one node execute the *startExample* function as follows:

```

1 (bob@localhost)3> diffie_hellman:startExample().
2 "Alice" (<0.50.0>): The shared private Key, exchanged with "Bob" is
   : 2
3 "Bob" (<0.51.0>): The shared private Key, exchanged with "Alice" is
   : 2
4 {<0.50.0>,<0.51.0>}

```

Executing the function *startExample* spawns a process that executes the *listenKeyExchange* function with 15 as the *private key* and "Alice" as the name and binds the returning pid to the value called Alice. Afterwards, *startExample* spawns another process that executes the *startKeyExchange* function with 23 for *P*, 5 for *G*, 6 for the *private key*, the pid of Alice and "Bob" as the name which binds the returning pid to the value called Bob. At the end the function returns the two pids of Alice and Bob. As we can see from the output, Alice has the pid <0.50.0> and the calculated private key, exchanged with Bob is 2. Bob has the pid <0.51.0> and has computed the same private key.

4.4.3 On Two Nodes

To run the application on two nodes open another shell, go to the *source* directory of the Erlang code and type:

```
$ erl -sname alice
```

Which will output (meta data depends of the executing system):

```
1 Erlang/OTP 17 [erts-6.3] [source] [64-bit] [smp:4:4] [async-threads
   :10] [hipe] [kernel-poll:false]
2
3 Eshell V6.3 (abort with ^G)
4 (alice@localhost)1>
```

Now another Erlang shell runs on the node *alice* where the source files shall be compiled, too. Switching to the Erlang shell on node *bob* and executing the function *startRemoteExample* with the atom *alice@localhost* (specifies on which node the actor shall be spawned that executes the *listenKeyExchange* function) as the transferred parameter as follows:

```
1 (bob@localhost)4> diffie_hellman:startRemoteExample(alice@localhost
   ).
```

Prints out the following:

```
1 {<9879.54.0>,<0.58.0>}
2 "Alice" (<9879.54.0>): The shared private Key, exchanged with "Bob"
   is: 2
3 "Bob" (<0.58.0>): The shared private Key, exchanged with "Alice" is
   : 2
```

All transferred parameters are the same with the difference that the *Alice* actor is now located on the node *alice*. The output is still printed out in the *bob* node, since the Erlang io system recognizes where the process is spawned from and sends all the output to it (Ericsson AB, 2015, ch. 4.3.4 on p. 104).

5 CLOSING REMARKS

A my_math Erlang Module

```
1 %%% @author      Florian Willich
2 %%% @copyright   The MIT License (MIT) Copyright (c) 2014
3 %%%              University of Applied Sciences, Berlin, Germany
4 %%%              Florian Willich
5 %%%              For more detailed information, please read the
6 %%%              licence.txt in the erlang root directory.
7 %%% @doc         This is my math module for mathematical
8 %%%              functions not provided by the erlang standard
9 %%%              library.
10 %%% @end
11 %%% Created 2015-06-06
12
13 -module(my_math).
14 -author("Florian Willich").
15 -export([pow/2]).
16
17 %%% @doc Returns the value of Base to the power of Exponent.
18 %%%      If Base and Exponent is 0 the function returns
19 %%%      {error, undefined_arithmetic_expression},
20 %%%      The motivation to implement this function was that
21 %%%      there is no erlang standard library pow function
22 %%%      returning an integer.
23 %%% @end
24 -spec pow(integer(), integer()) -> number() | {error, atom()}.
25 pow(0, 0) ->
26     {error, undefined_arithmetic_expression};
27
28 pow(Base, 0) ->
29     case Base < 0 of
30         true -> -1;
31         false -> 1
32     end;
33
34 pow(Base, Exponent) ->
35     case Exponent < 0 of
36         true -> 1 / pow(Base, -Exponent, 0);
37         false -> pow(Base, Exponent, 0)
38     end.
39
40 %%% @doc Returns the value of Base to the power of Exponent.
41 %%%      Acc should be 0 for initiating computation.
42 %%%      The motivation to implement this function was that
43 %%%      there is no erlang standard library pow function
44 %%%      returning an integer.
45 %%% @end
46 -spec pow(pos_integer(), non_neg_integer(), non_neg_integer()) ->
47     integer().
48 pow(_, 0, Acc) -> Acc;
49
50 pow(Base, Exponent, 0) ->
51     pow(Base, Exponent - 1, Base);
52
53 pow(Base, Exponent, Acc) ->
54     pow(Base, Exponent - 1, Acc * Base).
```

Acronyms

FIFO First In - First Out. 5

HPC High-Performance Computing. 5

IPC Inter-Process Communication. 3

ISO International Organization for Standardization. 4

JSON JavaScript Object Notation. 4

MPI Message-Passing Interface. 5, 6

MPHF Message-Passing Interface Forum. 5, 6

OS Operating System. 4, 7

OSI Open Systems Interconnection Reference Model. 4

pid Process Identifier. 7, 12

RPC Remote Procedure Calls. 3

RTP Real-Time Transport Protocol. 4

TCP/IP Transmission Control Protocol / Internet Protocol. 3, 4, 6

UDP User Datagram Protocol. 4

XML Extensible Markup Language. 4

References

- Armstrong, Joe. 1996. *Erlang - A survey of the language and its industrial applications*. In: *Proceedings of the symposium on industrial applications of Prolog (INAP96)*. <http://www.erlang.se/publications/inap96.ps>. The 9'th Exhibitions and Symposium on Industrial Applications of Prolog. 16-18, October 1996. Hino, Tokyo Japan. [Online. Accessed 5th May 2015].
- Däcker, Bjarne. 2000. *Concurrent functional programming for telecommunications: A case study of technology introduction*. <http://www.erlang.se/publications/bjarnelic.ps>. Licenciate Thesis, Department of Teleinformatics. TRITA-IT AVH 00:08, ISSN 1403-5286. Royal Institute of Technology, Stockholm, Sweden. [Online. Accessed 20th May 2015].
- Ericsson AB. 2015. *Erlang/OTP System Documentation 6.4*. <http://www.erlang.org/doc/pdf/otp-system-documentation.pdf>. [Online. Accessed 31th May 2015].
- Hebert, Fred. 2013. *Learn You Some Erlang for Great Good! : A Beginner's Guide*. <http://learnyousomeerlang.com/content>. No Starch Press. [Online. Accessed 20th May 2015].
- Message Passing Interface Forum. 2012. *MPI: A Message-Passing Interface Standard Version 3.0*. <http://www.mpi-forum.org/docs/mpi-3.0/mpi30-report.pdf>. [Online. Accessed 20th May 2015].
- Tanenbaum, Andrew S., & Steen, Maarten Van. 2007. *Distributed Systems: Principles and Paradigms (Second Edition)*. Pearson Prentice Hall. ISBN 0-13-239227-5.